

*(This problem is an interactive problem.)*A binary matrix means that all elements are `0` or `1`. For each **individual** row of the matrix, this row is sorted in non-decreasing order.Given a row-sorted binary matrix `binaryMatrix`, return leftmost column index(0-indexed) with at least a `1` in it. If such index doesn't exist, return `-1`.You can't access the **Binary Matrix directly**. You may only access the matrix using a `BinaryMatrix` interface:

- `BinaryMatrix.get(row, col)` returns the element of the matrix at index `(row, col)` (0-indexed).
- `BinaryMatrix.dimensions()` returns a list of 2 elements `[rows, cols]`, which means the matrix is `rows * cols`.

Submissions making more than `1000` calls to `BinaryMatrix.get` will be judged Wrong Answer. Also, any solutions that attempt to circumvent the judge will result in disqualification.For custom testing purposes you're given the binary matrix `mat` as input in the following four examples. You will not have access the binary matrix directly.**Example 1:**

0	0
1	1

Input: mat = [[0,0],[1,1]]

Output: 0

Example 2:

0	0
0	1

Input: mat = [[0,0],[0,1]]

Output: 1

Example 3:

0	0
0	0

Input: mat = [[0,0],[0,0]]

Output: -1

Example 4:

0	0	0	1
0	0	1	1
0	0	1	1

Input: mat = [[0,0],[0,1],[0,1]]

Output: 1

Constraints:

- `rows == mat.length`
- `cols == mat[0].length`
- `1 <= rows, cols <= 100`
- `mat[i][j]` is either `0` or `1`.
- `mat[i]` is sorted in a non-decreasing way.

Solution**Approach 1: Linear Search Each Row****Intuition**

This approach won't pass, but we'll use it as a starting point. Also, it might be helpful to you if you just needed an example of how to reply to the API, but don't want to see a complete solution yet!

The leftmost `1` is the `1` with the lowest column index.The problem can be broken down into finding the index of the first `1` in each row and then taking the minimum of those indexes.

The simplest way of doing this would be a linear search on each row.

Algorithm

```
Java Python3
1 class Solution:
2     def leftMostColumnWithOne(self, binaryMatrix: "BinaryMatrix") -> int:
3         rows, cols = binaryMatrix.dimensions()
4         smallest_index = cols
5         # Set pointers to the top-right corner.
6         current_row = 0
7         current_col = cols - 1
8         while current_row < rows:
9             for row in range(current_row, rows):
10                 if binaryMatrix.get(row, current_col) == 1:
11                     smallest_index = min(smallest_index, row)
12             current_col -= 1
13             if current_col == -1:
14                 return -1
15         return smallest_index
```

Complexity Analysis

If you ran this code, you would have gotten the following error.

You made too many calls to `BinaryMatrix.get()`.The maximum grid size is `100` by `100`, so it would contain `10000` cells. In the worst case, the linear search algorithm we implemented has to check every cell. With the problem description telling us that we can only make up to `1000` API calls, this clearly isn't going to work.Let `N` be the number of rows, and `M` be the number of columns.Time complexity: $O(N \cdot M)$ We don't know the time complexity of `binaryMatrix.get()` as its implementation isn't our concern. Therefore, we can assume it's $O(1)$.In the worst case, we are retrieving a value for each of the $N \cdot M$ cells. At $O(1)$ per operation, this gives a total of $O(N \cdot M)$.Space complexity: $O(1)$.

We are only using constant extra space.

Approach 2: Binary Search Each Row**Intuition**

This isn't the best approach, but it passes, and coding it up is a good opportunity to practice binary search.

When linear search is too slow, we should try to find a way to use binary search. If you're not familiar with binary search, check out this [Explore Card](#). We recommend doing the first couple of binary search questions to get familiar with the algorithm before coming back to this problem.Also, have a go at [First Bad Version](#). The only difference between that problem and this one is that instead of `0` and `1`, it uses `false` and `true`.Like we did with the linear search, we're going to apply binary search independently on each row. The target element we're searching for is the first `1` in the row.

The core part of a binary search algorithm is how it decides whether the target element is to the left or the right of the middle element. Let's figure this out by thinking through a couple of examples.

In the row below, we've determined that the middle element is a `0`. On what side must the target element (first `1`) be? The left, or the right? Don't forget, all the `0`'s are before all the `1`'s.

In this next row, the middle element is a `1`? What side must the target element be on? Could it also possibly be the `1` we just found?

For the first example, we can conclude that the target element (if it exists) must be to the **right** of the middle element. This is because we know that everything to the left of a `0` must also be a `0`.For the second example, we can conclude that the target element is either the middle element itself or it is some other `1` to the **left** of the middle element. We know that everything to the right of a `1` is also a `1`, but these can't possibly be further left than the one we just found.

In summary, if the middle element is:

- `0`, then the target must be to the **right**.
- `1`, then the target is either this element or to the **left**.

We can then put this together into an algorithm that finds the index of the target element (first `1`) in each row, and then returns the minimum of those indexes. Here is an animation of how that algorithm would look. The light grey numbers are the ones that we could *infer* without needing to make an API call. They are only there to help you understand.

API Calls: 0

As always in binary search, there are a couple more key implementation details we still need to deal with:

1. An even-length search space has two middles. Which do we choose?

2. The row might be all `0`'s.

Let's tackle these one at a time.

The first issue is the choice of middle, as otherwise, the search space might stop shrinking when it gets down to two elements. When the search space doesn't shrink, the algorithm does the exact same thing the next loop cycle, resulting in an infinite loop. Remember that when the search space only contains two elements, they are the ones pointed to by `lo` and `hi`. We, therefore, need to think about which cases would shrink the search space, and which would not.If we use the **lower-middle**:

- If it is a `0`, then we set `lo = mid + 1`. Because `hi == mid + 1`, this means that `lo < hi` (search space is of length zero).
- If it is a `1`, then we set `hi = mid`. Because `mid == lo`, this means that `hi == lo` (search space is of length one).

We are only using constant extra space.

Approach 2: Binary Search Each Row**Intuition**

This isn't the best approach, but it passes, and coding it up is a good opportunity to practice binary search.

When linear search is too slow, we should try to find a way to use binary search. If you're not familiar with binary search, check out this [Explore Card](#). We recommend doing the first couple of binary search questions to get familiar with the algorithm before coming back to this problem.Also, have a go at [First Bad Version](#). The only difference between that problem and this one is that instead of `0` and `1`, it uses `false` and `true`.Like we did with the linear search, we're going to apply binary search independently on each row. The target element we're searching for is the first `1` in the row.

The core part of a binary search algorithm is how it decides whether the target element is to the left or the right of the middle element. Let's figure this out by thinking through a couple of examples.

In the row below, we've determined that the middle element is a `0`. On what side must the target element (first `1`) be? The left, or the right? Don't forget, all the `0`'s are before all the `1`'s.

In this next row, the middle element is a `1`? What side must the target element be on? Could it also possibly be the `1` we just found?

For the first example, we can conclude that the target element (if it exists) must be to the **right** of the middle element. This is because we know that everything to the left of a `0` must also be a `0`.For the second example, we can conclude that the target element is either the middle element itself or it is some other `1` to the **left** of the middle element. We know that everything to the right of a `1` is also a `1`, but these can't possibly be further left than the one we just found.

In summary, if the middle element is:

- `0`, then the target must be to the **right**.
- `1`, then the target is either this element or to the **left**.

We can then put this together into an algorithm that finds the index of the target element (first `1`) in each row, and then returns the minimum of those indexes. Here is an animation of how that algorithm would look. The light grey numbers are the ones that we could *infer* without needing to make an API call. They